ROWAN BRAD QUNI-GUDZINAS

Autonomous AI systems researcher

Amsterdam, Netherlands · rowan.quni@qnfo.org · qnfo.org · ORCID 0009-0002-4317-5604 · LinkedIn

Earlier work is published under **Brad Gudzinas**.

SUMMARY

I study autonomous AI systems: how they are built, how they behave once they are running, how they fail, and how people can audit and govern them. The path here was long. I started in transportation network and GIS work in 2004, built network models at DVRPC and IBI Group, specified the national Tour-Based Model System at the U.S. Federal Highway Administration, ran data-driven decision support at AARP, and took on machine-learning engagements at Deloitte. Since 2024 I have run QNFO, an autonomous research system of 39 live agents, and published what it delivers, what it costs and where it fails. In 2026 I began turning that fleet into a product, Ultrastable (ultrastable.io): a development platform that chooses its own work, proves each change live and scores its own decisions. I direct it; coding agents write its code.

I am looking for work in autonomous and AI systems research and strategy, AI evaluation and assurance, AI governance and oversight, research program leadership for AI, or founder-in-residence and product strategy roles for autonomous systems.

EXPERIENCE

Founder and Director

Ultrastable (autonomous development platform, ultrastable.io) | 2026 to present

- Launching the system that runs QNFO as its own platform. It finds work in its own metrics and failures, hands each change to coding agents from several vendors, checks every deploy with live probes and rolls back automatically, and scores its decisions against predictions it logged first.
- My role is systems design and direction, not coding. I set the goals, standing directives and limits, and I judge how the platform behaves; the coding agents write the code.
- Wrote and revised the operating doctrine: what agents decide on their own, what needs a person, how every item is accounted for, and the limits no agent may cross. It went through five revisions in October 2026 and is published in QNFO/qnfo-workers.
- Drew the human-oversight line. Agents proceed on reversible work; only irreversible steps tied to a person's identity (money, legal or signing acts, the owner's own accounts) come to a person.
- Set the verification rules: no change counts until a live check proves it, every problem found is fixed, deployed and verified, and quality is judged by outcomes rather than proxy scores.
- Found that proxy signals (test pass rates, reviewer ratings, issue counts) can look healthy while the platform drifts. The drift showed only when decisions were scored against the predictions logged beforehand.
- Direct coding agents from several vendors (Claude Code, GitHub Copilot, OpenAI Codex) under one written operating agreement. Service levels are measured in minutes, every metric has a trigger, and AI spend caps steer model choice instead of stopping work.
- Evidence, re-checked by the running system and published live at ultrastable.io: 3,880 merged changes since August 2026; 774 of 1,059 scored decisions matched their logged prediction; 0 silent drops across 6,000+ closed issues.
- Stated limits: one installation so far (QNFO), hundreds of open issues against a target of zero, and decision scoring that the system grades itself.

Founder and Systems Researcher

QNFO (independent research on autonomous AI systems) | 2024 to present

- Designed and operate an autonomous research system on Cloudflare (39 live workers) that drafts, referees, revises, versions, publishes and corrects papers, with standing objectives, health checks and a failure ledger.
- Built QNFO as the first installation (customer zero) of Ultrastable, the platform listed above.
- Reported how the system behaves in operation: Operating the Quniverse Fleet sets out its objectives, successes, failures and roadmap.
- Designed oversight methods for AI-assisted work: The Universal Ignorance Audit and Epistemic Legibility in AI-Assisted Science.
- Proposed JPCUB, a whole-system cost measure for computing (joules per correct answer), with a measurement protocol and an extension to stochastic and agentic LLM inference.
- Publish an open library of preprints with DOIs (papers.qnfo.org) and a knowledge graph that other tools can query by API.

Executive Director and Founder

Empowering Change (nonprofit) | San Francisco, CA | 2023 to 2024

- Took the nonprofit from incorporation to public launch.
- Led an LLM legal-navigation tool for self-represented litigants: document drafting, procedural guidance and plain-language translation of legal terms.

Platform Data Product Manager

iManage | 2022 to 2023

Senior Product Manager

Epsilon (Publicis Groupe) | 2021 to 2022

- Managed data products built on predictive machine-learning models.

Product Manager and Senior Methods Advisor

AARP Public Policy Institute | Washington, DC | 2016 to 2021

- Led the AARP Livability Index, a national decision-support system: 50+ data sources across 7 domains, scoring U.S. neighborhoods nationwide, across multiple public releases.
- Co-authored the Livability Index 2018 report and AARP policy work on housing affordability, including the 2016 analysis of Low-Income Housing Tax Credits.
- The Index and its method are cited in 20+ academic and policy studies.

Senior Data Scientist

The Advisory Board Company | 2016

Senior Associate, Data Analytics

Deloitte Consulting LLP | Washington, DC | 2015 to 2016

- Delivered analytics and machine-learning engagements for federal and public-sector clients.
- Co-authored a 2016 Transportation Research Board paper on updating the National Household Travel Survey with American Community Survey data.

Data Analyst and Research Manager

U.S. Department of Transportation, Federal Highway Administration | Washington, DC | 2011 to 2015

- Specified the national Tour-Based Model System, a simulation of long-distance passenger travel across the United States, and directed the contractor team that built it.

- Co-directed the $10M U.S. DOT National Household Travel Survey.
- Managed a $1.5M+ federal research portfolio as a certified Contracting Officer's Representative: technical oversight, budgets, milestones and acceptance of contractor deliverables.
- Presented national long-distance travel data needs at the 2011 TRB National Household Travel Survey workshop.

Earlier roles: transportation networks, systems planning and GIS

Role	Organization	Period
Consultant: built and ran transportation network models	IBI Group	2008 to 2011
Systems Planner: built and ran transportation network models	Delaware Valley Regional Planning Commission	2005 to 2008
GIS Analyst	Cross County Connection TMA	2004 to 2005

SELECTED WORK

1. Operating the Quniverse Fleet: Objectives, Successes, Failures, and Roadmap of an Autonomous Cloud Research System (2026)
2. The Universal Ignorance Audit (2026)
3. Epistemic Legibility in AI-Assisted Science (2026)
4. Joules-per-Solution for Stochastic and Agentic Inference (2026)
5. Tour-Based National Model System, FHWA (2014, as Brad Gudzinas)
6. AARP Livability Index 2018 report (co-author, as Brad Gudzinas)

Every QNFO paper is prepared with an AI-assisted research pipeline; the author is responsible for its content. Full library: papers.qnfo.org. All records: Zenodo.

EDUCATION

- **M.S., City and Regional Studies**, Rutgers University, Edward J. Bloustein School of Planning and Public Policy (2008)
- **B.A., Urban Studies and Community Planning**, Rutgers University-Camden (2004)

CERTIFICATIONS

- Project Management Professional (PMP), Project Management Institute (2016)
- Contracting Officer's Representative (COR), Defense Acquisition University (2011)
- American Institute of Certified Planners (AICP), American Planning Association (2009)

SKILLS

Autonomous systems direction: autonomous systems governance and doctrine design · human oversight and approval-boundary design for AI agents · verification and evaluation design for AI agents · failure-mode and risk analysis · multi-vendor AI agent orchestration and oversight · operating metrics, service levels and AI cost governance · writing machine-checkable specifications and requirements · product and platform strategy for autonomous systems · oversight and audit methods for AI-assisted work

Systems and decisions: systems and network modeling · travel-demand and simulation modeling · statistical modeling and applied machine learning · data-driven decision support · GIS · research program direction and management · federal contract oversight

Tools I direct and work with: Python, SQL, Cloudflare Workers, GIS (ArcGIS, QGIS), LLM agent frameworks.